Class 1: What is AI?

Principles of ML & Deep Learning - SAILea

Anisha Raghu and Ananya Raghu

Icebreakers

- Name
- Grade
- A fun fact about yourself :)
- How much do you know about AI? (it's ok if you do not have much past experience)

Machine Learning Models

- Decision Trees
- Random Forest
- Linear & Logistic Regression
- K-nearest neighbors
- Etc…

What is AI?

- Based on patterns
- Process large amounts of data and make predictions from it
- Traditional programming: does step by step commands - like a recipe using ingredients
- AI: Give computer various ingredients, and various outputs, and have it build or *learn* the recipe

Supervised, Unsupervised Learning, Reinforcement

- In simple terms, *supervised learning* is the process of using **labeled** examples to predict unlabeled examples

- *Unsupervised learning*: uses unlabeled data, discovers "structure" or underlying patterns in data

- *Machine Learning:* Process of learning to pick actions based on rewards and punishments from previous choices

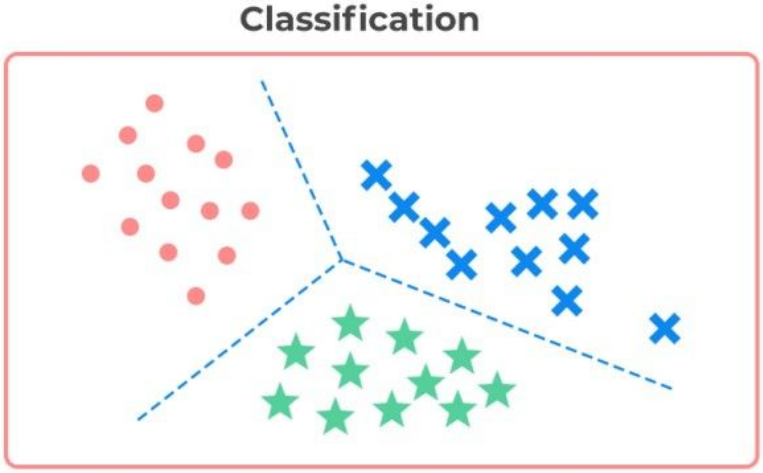

Supervised learning

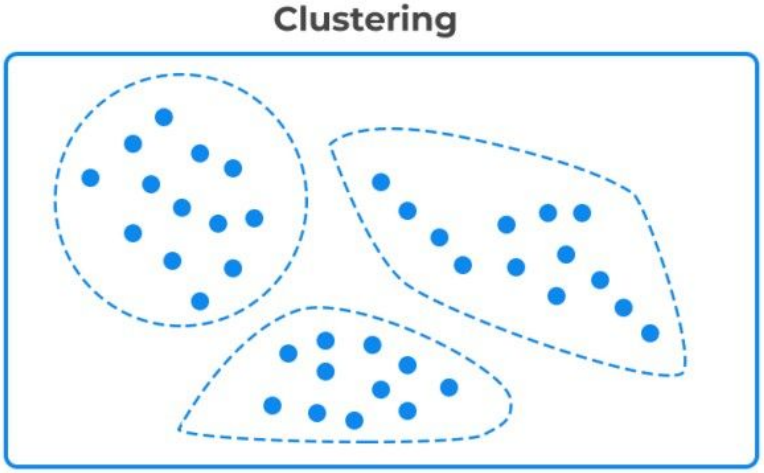

Unsupervised learning

Regression vs Classification

- Regression helps predict CONTINUOUS values
- Classification simply classifies or predicts DISCRETE class labels

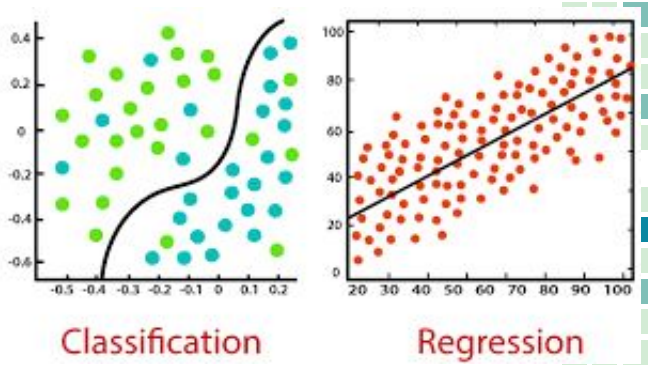

Source: https://www.javatpoint.com/regression-vs-classification-in-machine-learning

Practice question: I want to predict the temperature tomorrow in San Francisco. Would this be a classification or regression problem? What if I wanted to ask if it will rain or be sunny (assuming these are the only possible scenarios)?

Data

- While implementing machine learning, we generally split our data into three sets:
  - 1) **Training Data** - data used to fit the model
  - 2) **Validation Data** - Used to evaluate model performance at each step (epoch) of training. Note: model "sees" this, but does not "learn" directly from this
    - Example: I am designing a model - I would use validation dataset to decide hyperparameters of the model → for example, how many layers should be included?
  - 3) **Testing Data** - Used after training has completed to observe final performance of model
  -

Example Dataset for ML

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883

Each row can be thought of as a point in n-dimensional space, where n is the number of features!

Metrics

- After training our model, how can we evaluate its performance?
- Basic metric: **accuracy**
  - (number of correct decisions)/total decisions

Continued - Confusion Matrices

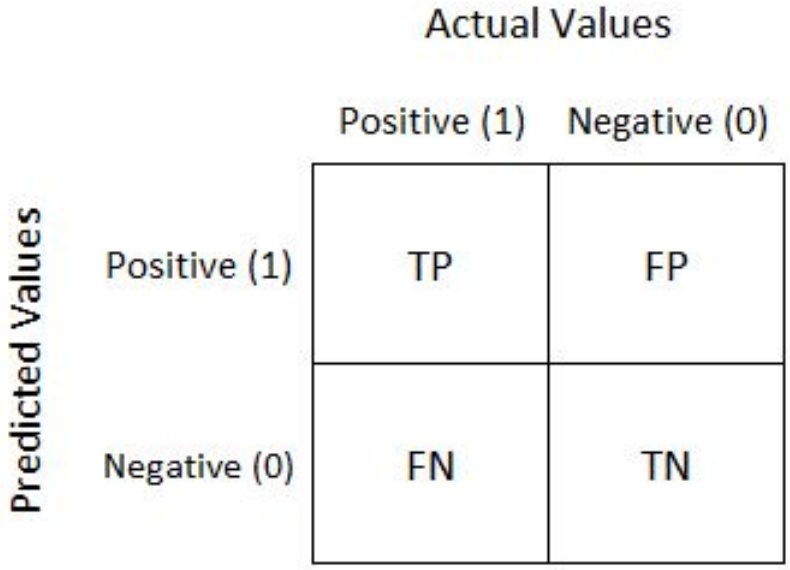

- Way to visualize the performance of an algorithm
- Summarizes performance of a classification algorithm

Source of Image:
https://towardsdatascience.com/understanding-confusion-matrix-a9ad42dcfd62

Overfitting versus Underfitting

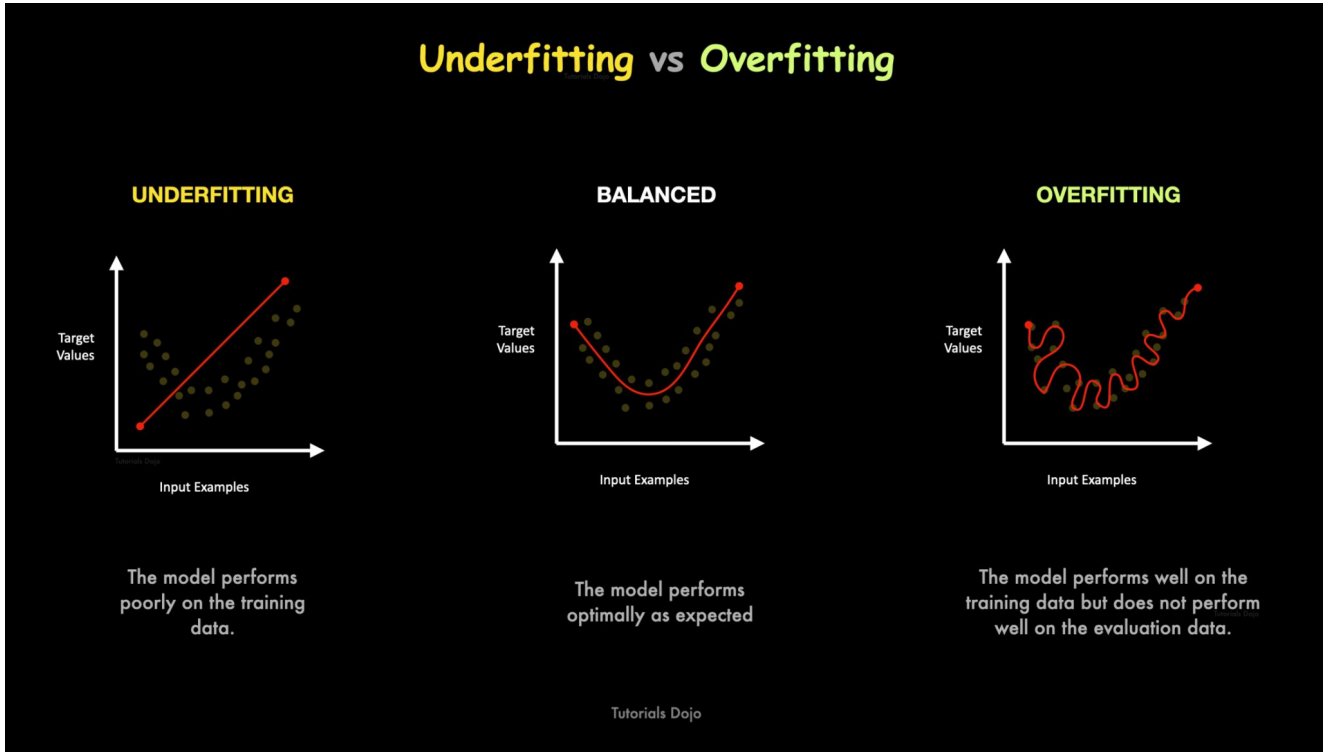

Source: https://www.linkedin.com/pulse/underfitting-vs-overfitting-simplified-jon-bonso

Feature Importance

- Indicates how much each feature contributes to the overall prediction
- **Permutation Feature Importance:** Permute one feature in the dataset and see how it affects the prediction.

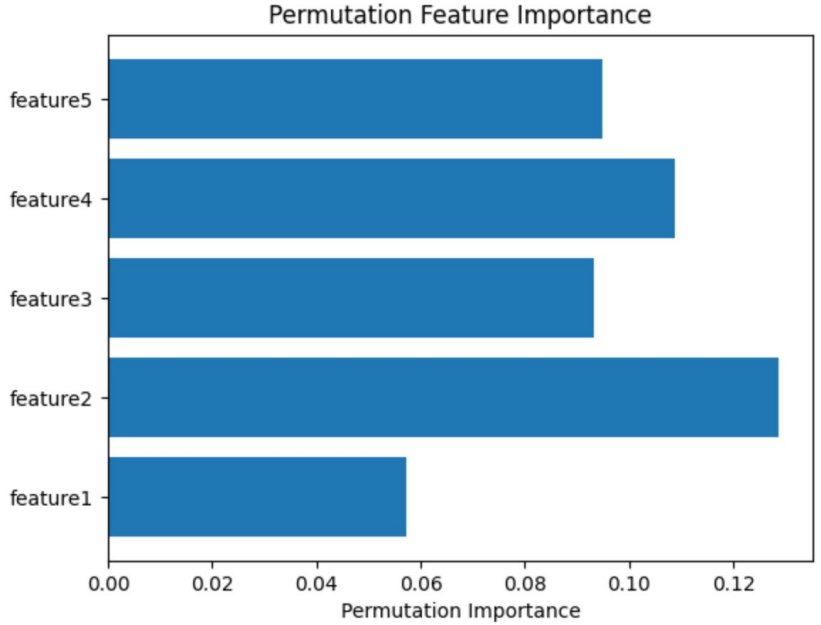

Source:https://towardsdatascience.com/feature-importance-in-machine-learning-explained-443e35b1b284

Decision Tree

- A decision tree is a supervised learning algorithm
- Used for both classification and regression tasks.
- Tree structure, which consists of a root node, branches, internal nodes and leaf nodes.

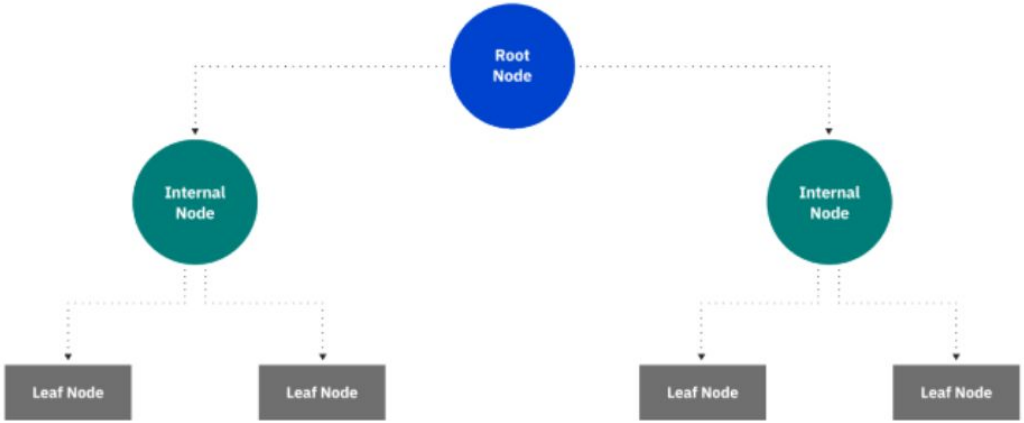

Decision Tree Explained

- How does a decision tree decide what is the condition at each node?
- A commonly used method is via the **gini** calculation, which measures "impurity" in a distribution:
  - $Gini = 1 - \sum_{i=1}^{n}(p_i)^2$ Where p_i is the probability of an object being classified to a particular class

Example: Gini Calculation

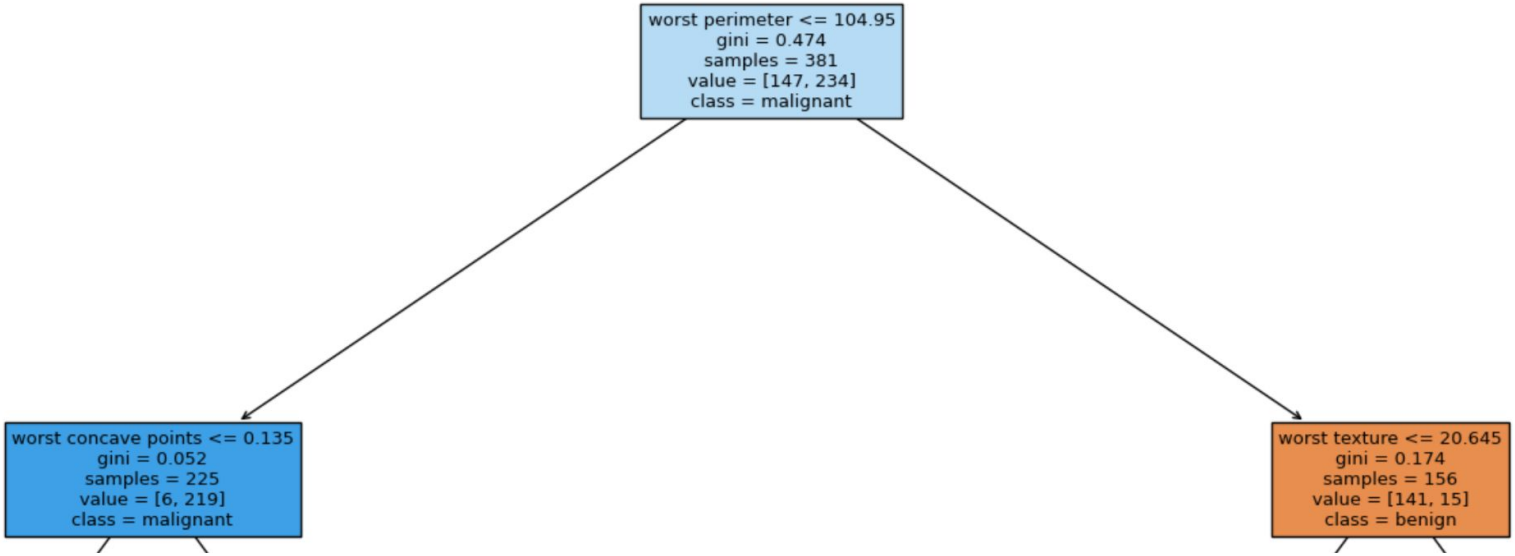

In the second layer, the gini value for the blue box is 0.052 because it is 1 - (6/225)^2 - (219/225)^2. This can be thought of as a collection of 6 blue marbles, 219 red marbles, which is close to being very pure.

How do Decision Trees use the Gini Index?

- First, the decision tree will choose a split function.
- Based on split function, finds probability of going to left node or going to right node (P_R , P_L)
- Then, can compute gini for left node and gini for right node
- The decision will try to pick the split function that minimizes P_L*$Gini_L$ + P_R*$Gini_R$

Any questions?